

SQL 1 : Basic Statements

Gary KL Tam

Department of Computer Science
Swansea University

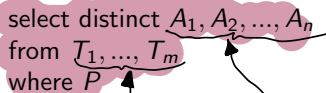
- Structured query language (**SQL**) is a user-friendly language for specifying relational algebra queries. It is supported by all the major database systems.

SQL provides:

- Data Manipulation Language (DML)
 - retrieve, insert and modify database contents
- Data Definition Language (DDL)
 - add and delete database objects
- Data Control Language (DCL)
 - configure security access
- In this lecture, we will learn how to rewrite algebra operators in SQL (DML).

Syntax of an SQL Statement

select distinct A_1, A_2, \dots, A_n
from T_1, \dots, T_m
where P




where T_1, \dots, T_m are tables, A_1, \dots, A_n are attributes, and P is a predicate. The statement returns a table, and corresponds to the following relational algebra query:

$$\Pi_{A_1, \dots, A_n}(\sigma_P(T_1 \times \dots \times T_m))$$

Selection σ

corresponds to



PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500
<i>p6</i>	Frank	CS	full	9000

```
select *
from PROF
where rank = 'asst'
```

$$\sigma_{\text{rank} = \text{"asst"}}(\text{PROF})$$

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

```
select *
from PROF
where not(rank = 'asst' and dept = 'EE')
```

$\sigma_{\neg(\text{rank} = \text{"asst"} \wedge \text{dept} = \text{"EE"})}(\text{PROF}) \rightsquigarrow \text{Sal Statement}$

Selection Predicate

```
select *  
from T  
where P
```

In P , you can specify the standard comparisons and logic operators:

- $=, <>, <, <=, >, >=$ *boolean expression.*
- Connect multiple comparisons with: **AND, OR, NOT.**

Projection Π

corresponds to

select distinct A_1, \dots, A_n
from T

\rightarrow Table

$\Pi_{A_1, \dots, A_n}(T)$

distinct
attribute

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

select distinct dept, rank
from PROF

$\Pi_{\text{dept, rank}}(\text{PROF})$

Note

The keyword **distinct** removes all duplicate rows in the output. Omitting the keyword keeps all duplicates. See the next slide.

PROF

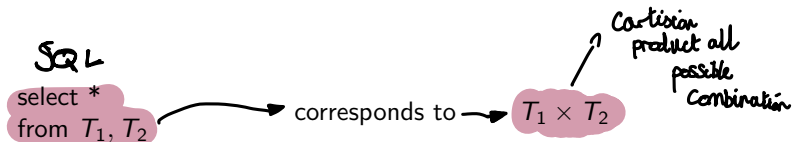
pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

“select dept, rank from PROF” returns:

dept	rank
CS	asst
EE	asso
CS	full
EE	asst
EE	asso
CS	full

This duplicate-retaining feature is useful for aggregate queries as we will discuss later in the course.

Cartesian Product \times



PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500

TEACH

pid	cid	year
<i>p1</i>	<i>c1</i>	2011
<i>p2</i>	<i>c2</i>	2012
<i>p1</i>	<i>c2</i>	2012

```
select *
from PROF, TEACH
```

~~PROF~~  ~~TEACH~~

Putting Multiple Operators Together

PROF

pid	name	dept	rank	sal
<i>p1</i>	Adam	CS	asst	6000
<i>p2</i>	Bob	EE	asso	8000
<i>p3</i>	Calvin	CS	full	10000
<i>p4</i>	Dorothy	EE	asst	5000
<i>p5</i>	Emily	EE	asso	8500

TEACH

pid	cid	year
<i>p1</i>	<i>c1</i>	2011
<i>p2</i>	<i>c2</i>	2012
<i>p1</i>	<i>c2</i>	2012

```
select distinct dept  
from PROF, TEACH  
where PROF.pid = TEACH.pid
```

```
 $\Pi_{\text{dept}}(\sigma_{\text{PROF.pid} = \text{TEACH.pid}}(\text{PROF} \times \text{TEACH}))$ 
```

Rename ρ

```
select ...  
from T as S  
where ...
```

corresponds to

$\dots\rho_S(T)\dots$

Come back
later

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500

TEACH

pid	cid	year
p1	c1	2011
p2	c2	2012
p1	c2	2012

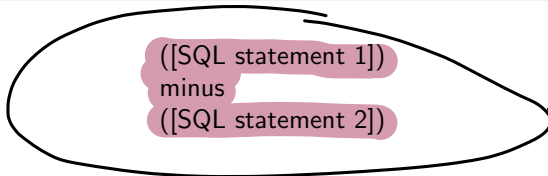
select distinct dept

from PROF as A, TEACH as B

where A.pid = B.pid

$$\Pi_{\text{dept}}(\sigma_{A.\text{pid} = B.\text{pid}}(\rho_A(\text{PROF}) \times \rho_B(\text{TEACH})))$$

Set Difference —



→ minus
will automatically
remove all
duplicates

↓
as it's a Set
theory because
Set word 'minus'

corresponds to

$$T_1 - T_2$$

where T_1 (T_2) is the table returned by SQL statement 1 (2).

Note

- T_1 and T_2 need to have the same schema.
- Duplicates in T_1 and T_2 will first be removed before performing the set difference.
- In some systems (e.g., SQL server from Microsoft), the set difference operator is named "except", instead of "minus".

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

(select rank from PROF)

minus

(select rank from PROF where dept = 'CS')

$$\Pi_{\text{rank}}(\text{PROF}) - \Pi_{\text{rank}}(\sigma_{\text{dept} = \text{"CS"}}(\text{PROF}))$$

Set Union \cup

([SQL statement 1])

union

([SQL statement 2])

→ Set theory
duplicates will
be removed.

corresponds to

$$\underline{\underline{T_1 \cup T_2}}$$

where T_1 (T_2) is the table returned by SQL statement 1 (2).

Note

- T_1 and T_2 need to have the same schema.
- Duplicates in T_1 and T_2 will first be removed before performing the set union.

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

(select * from PROF where sal <= 6000)

union

(select * from PROF where sal >= 9000)

$$\sigma_{sal \leq 6000}(\text{PROF}) \cup \sigma_{sal \geq 9000}(\text{PROF})$$

We have shown how to rewrite the 6 fundamental algebra operators in SQL. How about the extended operators \leftarrow , \cap , \bowtie and \div ? As we will see next, there is an explicit statement only for \cap . Nevertheless, as \cap and \bowtie can be implemented using the 6 fundamental operators, they can also be written in SQL using the statements introduced earlier. We will, however, ignore \leftarrow from our discussion (this operator is the least useful one, anyway).

Set Intersection \cap

```
([SQL statement 1])  
intersect  
([SQL statement 2])
```

corresponds to

$$T_1 \cap T_2$$

where T_1 (T_2) is the table returned by SQL statement 1 (2).

Note

- T_1 and T_2 need to have the same schema.
- Duplicates in T_1 and T_2 will first be removed before performing the set union.

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500
p6	Frank	CS	full	9000

(select * from PROF where sal >= 6000)
intersect
(select * from PROF where dept = 'CS')

$\sigma_{sal \geq 6000}(\text{PROF}) \cap \sigma_{dept = \text{"CS"}}(\text{PROF})$

Natural Join

PROF

pid	name	dept	rank	sal
p1	Adam	CS	asst	6000
p2	Bob	EE	asso	8000
p3	Calvin	CS	full	10000
p4	Dorothy	EE	asst	5000
p5	Emily	EE	asso	8500

TEACH

pid	cid	year
p1	c1	2011
p2	c2	2012
p1	c2	2012

select distinct PROF.pid, name, dept, rank, sal, cid, year
from PROF, TEACH
where PROF.pid = TEACH.pid

Can't explicitly write natural join
as not excepted in SQL

$\Pi_{\text{PROF.pid, name, dept, rank, sal, cid, year}}(\sigma_{\text{PROF.pid}=\text{TEACH.pid}}(\text{PROF} \times \text{TEACH}))$

$=$
 $\text{PROF} \bowtie \text{TEACH}$

Division



T_1

pid	cid
p1	c1
p1	c2
p1	c3
p2	c2
p2	c3
p3	c1
p4	c1
p4	c2
p4	c3

T_2

cid
c1
c2
c3

(select pid from T_1)

minus

select pid from (

(select * from (select pid from T_1), T_2)

minus

(select * from T_1))

Note

Notice how an SQL statement can be nested in a from clause.

$$\Pi_{S_1-S_2}(T_1) - \Pi_{S_1-S_2}(\Pi_{S_1-S_2}(T_1) \times T_2 - T_1) = T_1 \div T_2$$

Funny stuffs

```
SELECT * FROM politicians WHERE clue > 0;
```

```
SQL> select intelligence_level from developer;  
select intelligence_level from developer  
      *  
ERROR at line 1:  
ORA-00904: "INTELLIGENCE_LEVEL": invalid identifier
```

http://www.orafaq.com/wiki/Fun_stuff