

Concurrency

Gary KL Tam

Department of Computer Science
Swansea University

This lecture

Overview

- Concurrency Protocols
- Involves two parts:
 - Locking Protocol
 - Deadlock Detection and Handling

Concurrency Control - Review

- We already learned that every schedule executed by the database must be serializable.
- We discuss one important type of schedule:
conflict serialisable schedule
- Important questions
 - How do we test a schedule is not conflict serialisable?
 - Detecting cycle in Precedence Graph (DONE)
 - How do we construct conflict serialisable schedules?
 - focus of this lecture

Motivation

- Assume that the database receives several transactions now, and it needs to find a serializable schedule.

- Two possible ways:

- **Analyze** the statements of all transactions before execution. Rearrange operations.

- But the analysis time may be **too long**.

*Our mum/ super nan
use their experiences
to prepare muffins!
(plan in advance)*

- **Run** whatever statement that comes next immediately.

DBMS does this!

- Fast but needs to apply some “execution rules”.
- ==> **Concurrency protocols**.

This lecture


Overview

- Concurrency Protocols
- Involves two parts:
 - Locking Protocol
 - Deadlock Detection and Handling

Lock protocol – an Idea

- Ideas: transactions can proceed only after the necessary locks are obtained.
- If a *lock cannot be granted*, the requesting transaction must **wait** until other transactions release locks.

Some team work scenarios

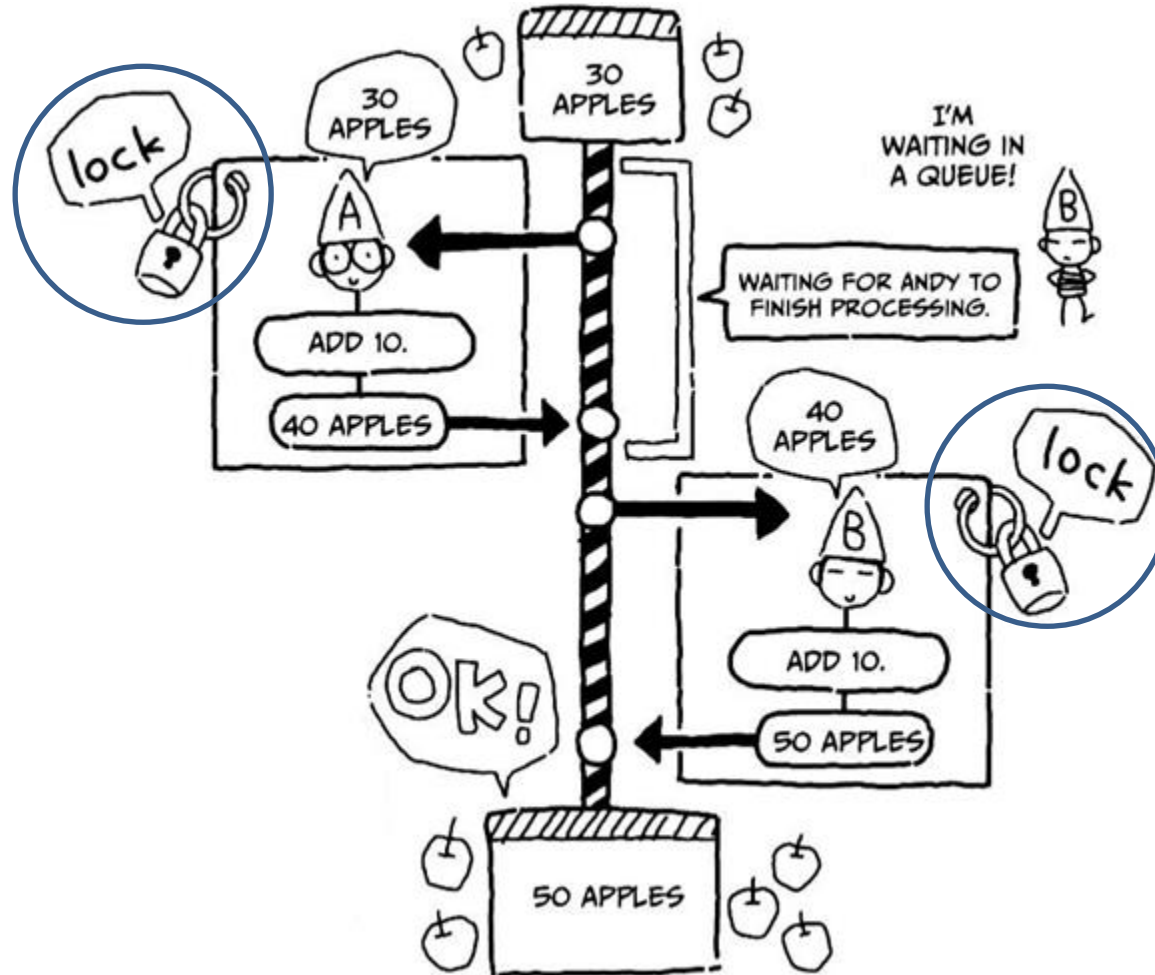


SVN: simple lock-modify-unlock
SVN is more advanced than this.

[Apache Subversion \(SVN\)](http://www.apache.org/licenses/LICENSE-2.0)

Lock protocols - inefficiency

- Locking the whole process leads to a serial schedule



- No concurrency....

Simple Browse / Check Balance

TESCO Every little helps

Sign in Store locator Contact us Help Rate this page

Search in Groceries for Keywords Search

Fresh food & groceries

Tesco direct

F&F Clothing

Clubcard

Tesco Bank

F&F

Online & In-Store

25% OFF

COATS & KNITWEAR

Clubcard Boost

DON'T FORGET TO BOOST YOUR VOUCHERS WITH F&F

Selected lines only. Discount automatically applied at checkout.

Shop F&F Clothing >

Women >

Men >

Kids >

School Uniform >

Real Food

Wine by the case

Discover blinkbox

Phone Shop

Health & Wellbeing

Fuel Save

Pri

- Should I lock all the resources (table/database)?



▶	On Screen Balance	Cash with Receipt	◀
▶	Postal Statement	Printed Balance	◀
▶	Cash without Receipt	Mini Statement	◀

Locking Protocol - Better idea

- **Two kinds of Locks:**
- Shared (S) lock on an item enables a transaction to **read** the item only. **Allow others to read.**
- Exclusive (X) lock on a data item enables a transaction to **read** and **write** to the item. **Exclude others from read/write.**
- A transaction obtains a lock only if the lock is **compatible** with those **already** on the data item.


Summary:

Shared (S) lock – Read Only

Exclusive (X) lock – Read and Write



Allow **other** transactions to obtain the lock?



Lock **already** acquired by one transaction

Read & Write

Read Only



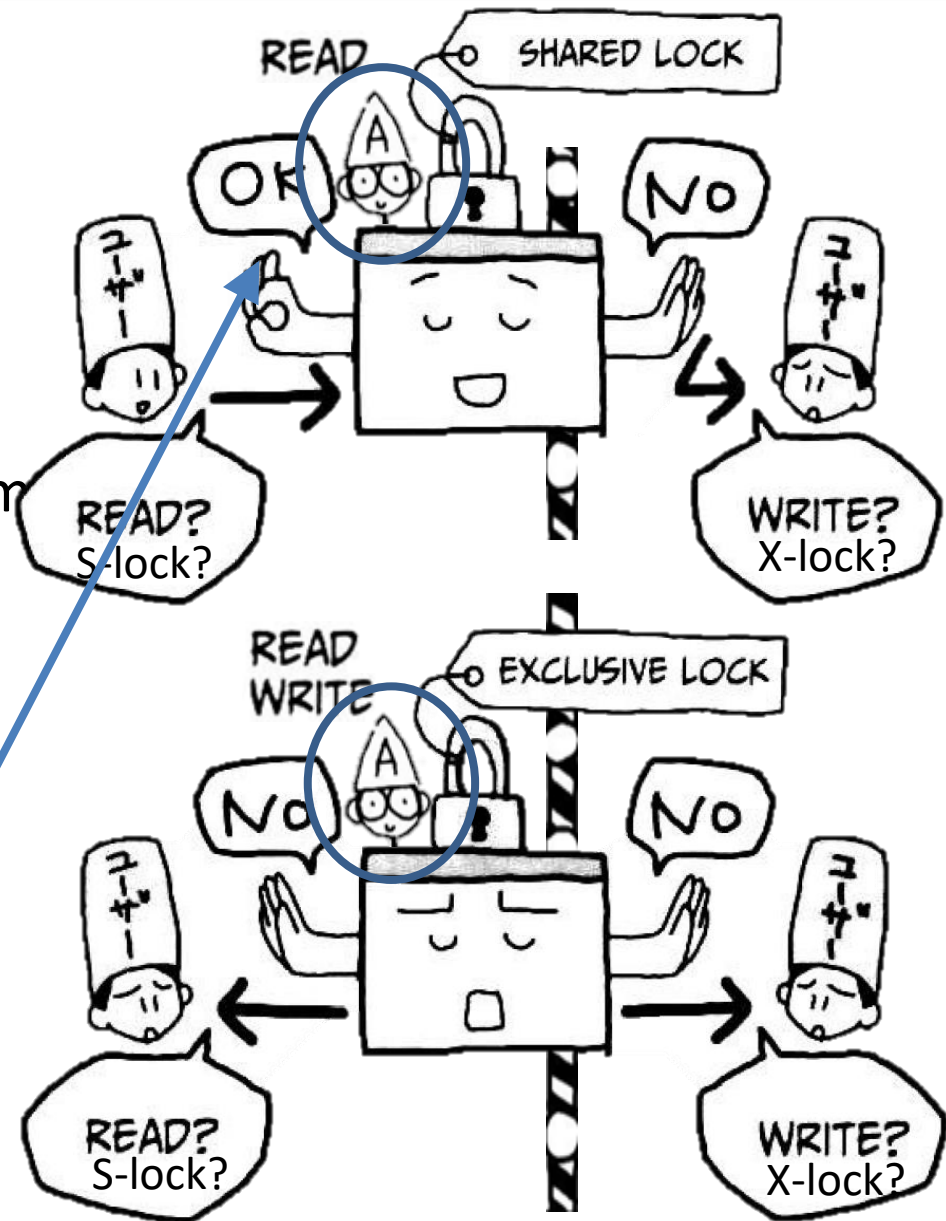
	Read & Write	Read Only
Read & Write	X	False
Read Only	False	True

Lock compatibility on the **same resource**

Locking Protocol - Better idea

- **Two kinds of Locks:**
- Shared (S) lock on an item enables a transaction to **read** the item only. **Allow others to read.**
- Exclusive (X) lock on a data item enables a transaction to **read** and **write** to the item. **Exclude others from read/write.**
- A transaction obtains a lock only if the lock is **compatible** with those **already** on the data item.

	X	S
X	False	False
S	False	True



More Locking Problems

- **Simple locking won't** allow us to serialise all schedules.
- For example:

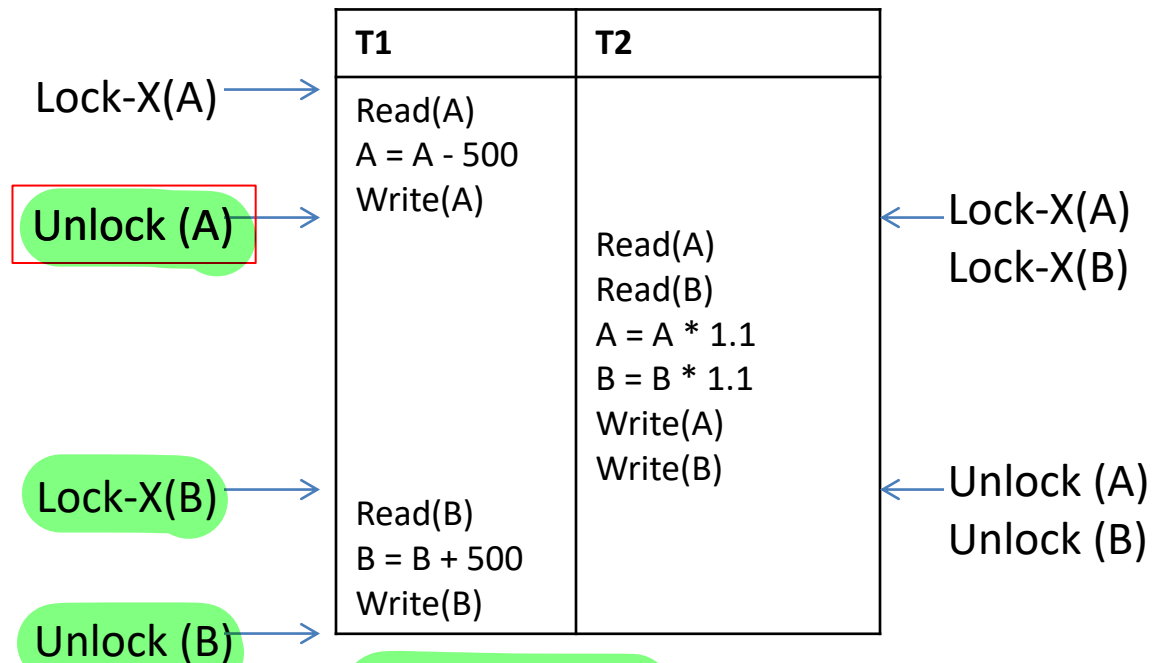
A:1000, B:2000

T1	T2
Read(A) A = A - 500 Write(A) Read(B) B = B + 500 Write(B)	Read(A) Read(B) A = A * 1.1 B = B * 1.1 Write(A) Write(B)

A:550, B:2750 (expected)

Serial schedule

A:1000, B:2000



A:550, B:2700



Two-Phase Locking

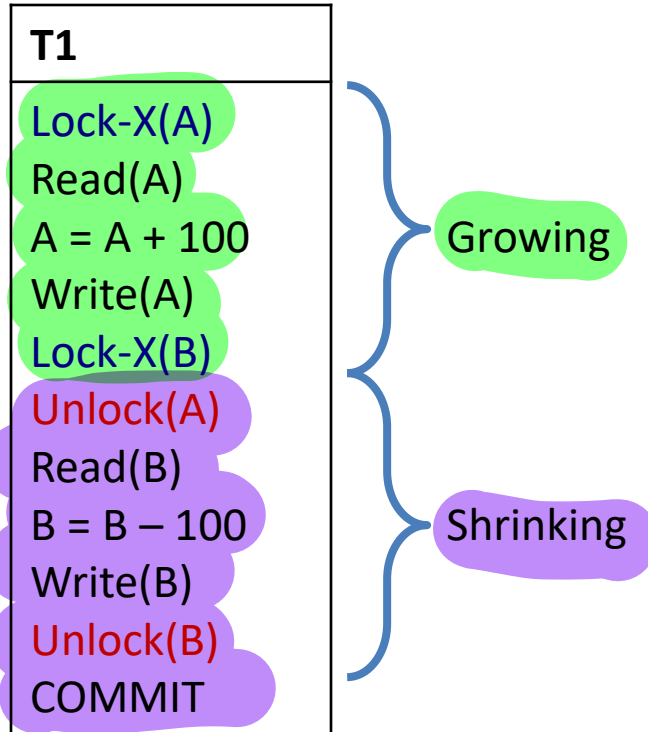
- A transaction follows two-phase locking protocol (2PL) if **all locking operations precede all unlocking operations**
 - Once a transaction releases a lock, it cannot apply for any lock in the future
- Two phases:
 - **Growing** phase where locks are acquired
 - **Shrinking** phase where locks are released

Serialisability Theorem

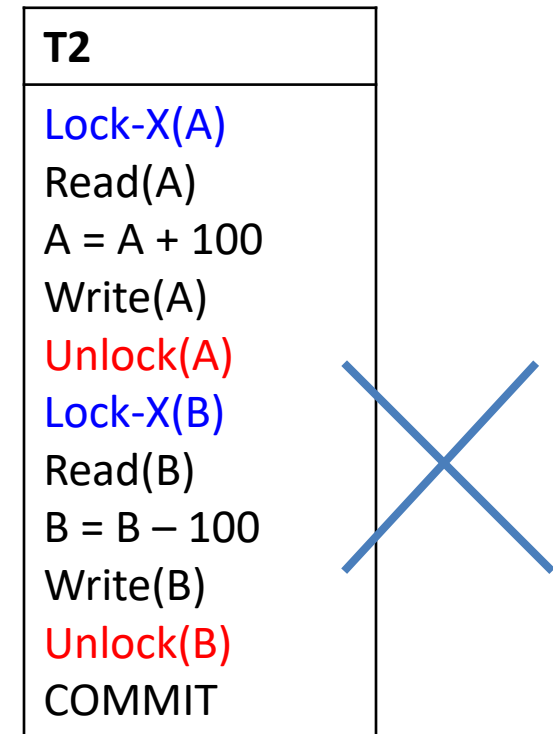
Any schedule of **two-phase locking** transactions is **conflict serialisable**

Two-Phase Locking Example

- T1 follows 2PL protocol
 - All locks in T1 are acquired **before** any are released
 - This happens even if the resource is no longer used

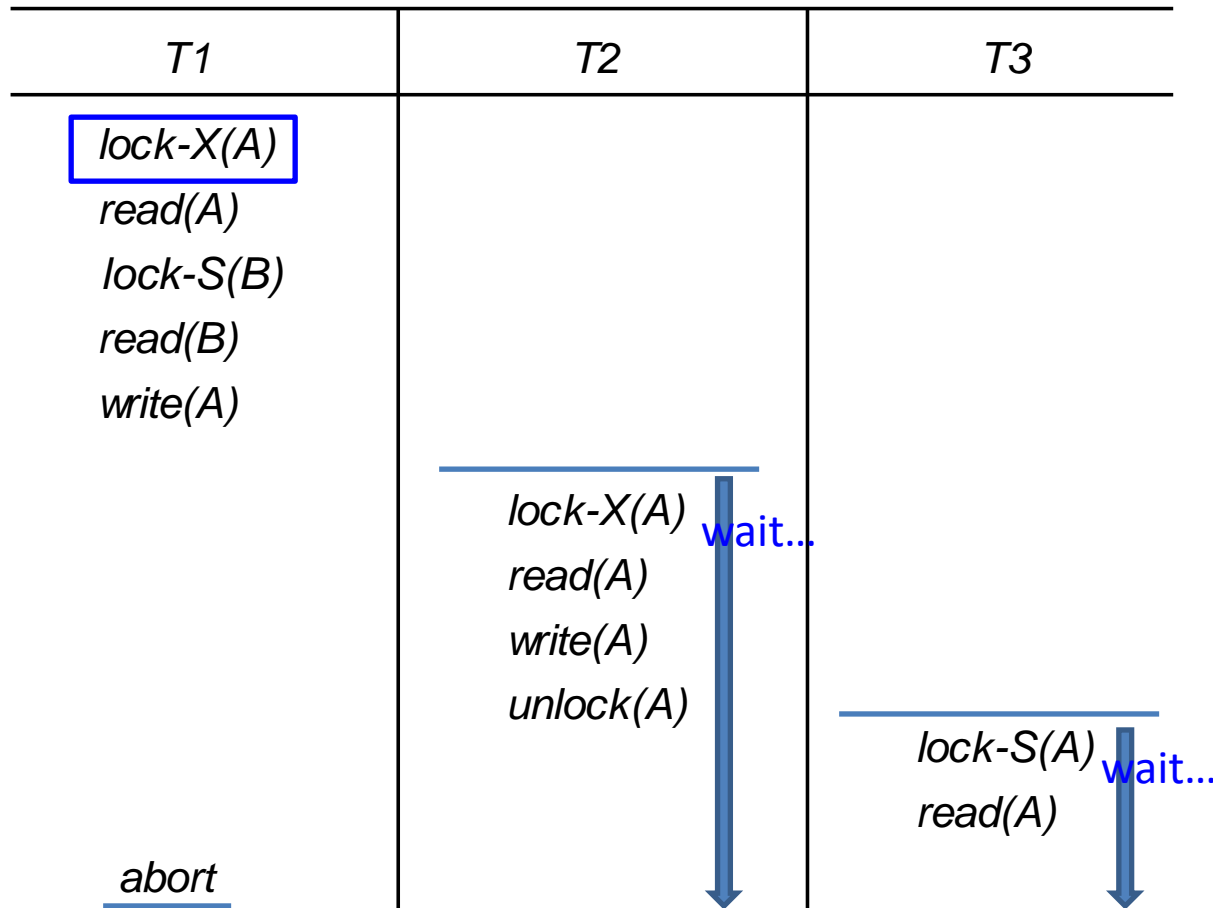


- T2 **does not follow 2PL**
 - Releases a lock on A, which is no longer needed, before acquiring on B



Strict Two-Phase Locking (**Strict** 2PL)

- Further requires that all locks be held by a transaction **until it commits/aborts**.



Easier to implement
by DBMS

- We now discuss strict 2PL in action.

Two-Phase Locking – Muddiest Points

- When to lock?
 - Tuple **retrieve** operation: Lock-S
 - Tuple **update** operation (i.e. insert, delete, update): Lock-X
- When to unlock?
 - 2PL locking: can unlock locks (that no longer in use) after all essential locks are acquired.
 - **Strict** 2PL Locking: unlock only at the end of transaction.
- Lock promotion in growing phase:
 - If a transaction T
 - holds a Lock-S on the tuple A, and
 - wishes to update tuple A
 - T must promote/upgrade Lock-S to Lock-X (according to lock compatibility).

Strict 2PL - Lost Update

- Lost Update

Lock	X	S
X	False	False
S	False	True

T1	T2
Read(X) $X = X + 10$	
	Read(X) $X = X + 10$
Write(X)	
	Write(X) COMMIT
COMMIT	



Lock-S(X) →

upgrade
Lock-X(X) →

T1	T2
Read(X) $X = X - 5$	
Write(X) WAIT WAIT ...	Read(X) $X = X + 5$ Write(X) WAIT WAIT ...

← Lock-S(X)

← upgrade
Lock-X(X)

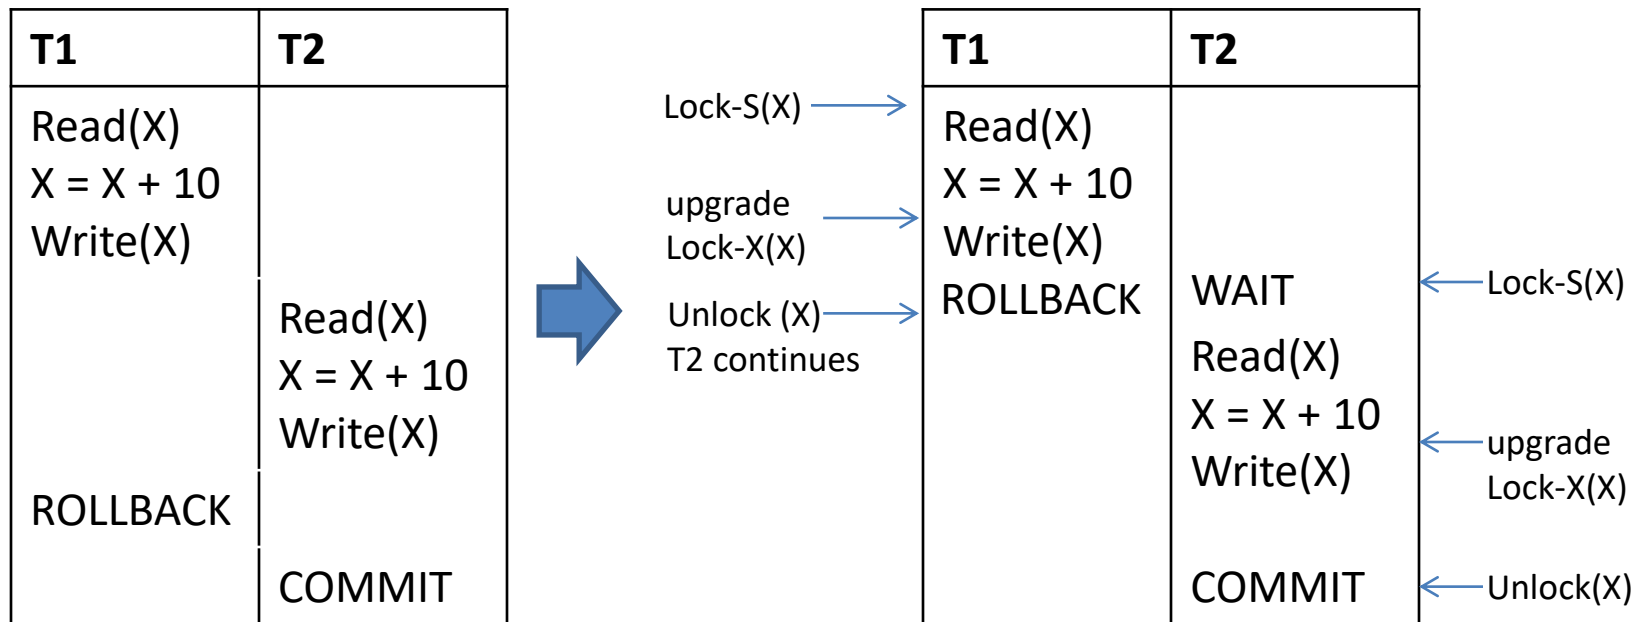
Shared lock on resource X from a different thread

Strict 2PL prevents Lost Update, and both not allow to continue.
That's important: No inconsistency in database. Good!

Strict 2PL - Uncommitted Update

Lock	X	S
X	False	False
S	False	True

- Uncommitted Update

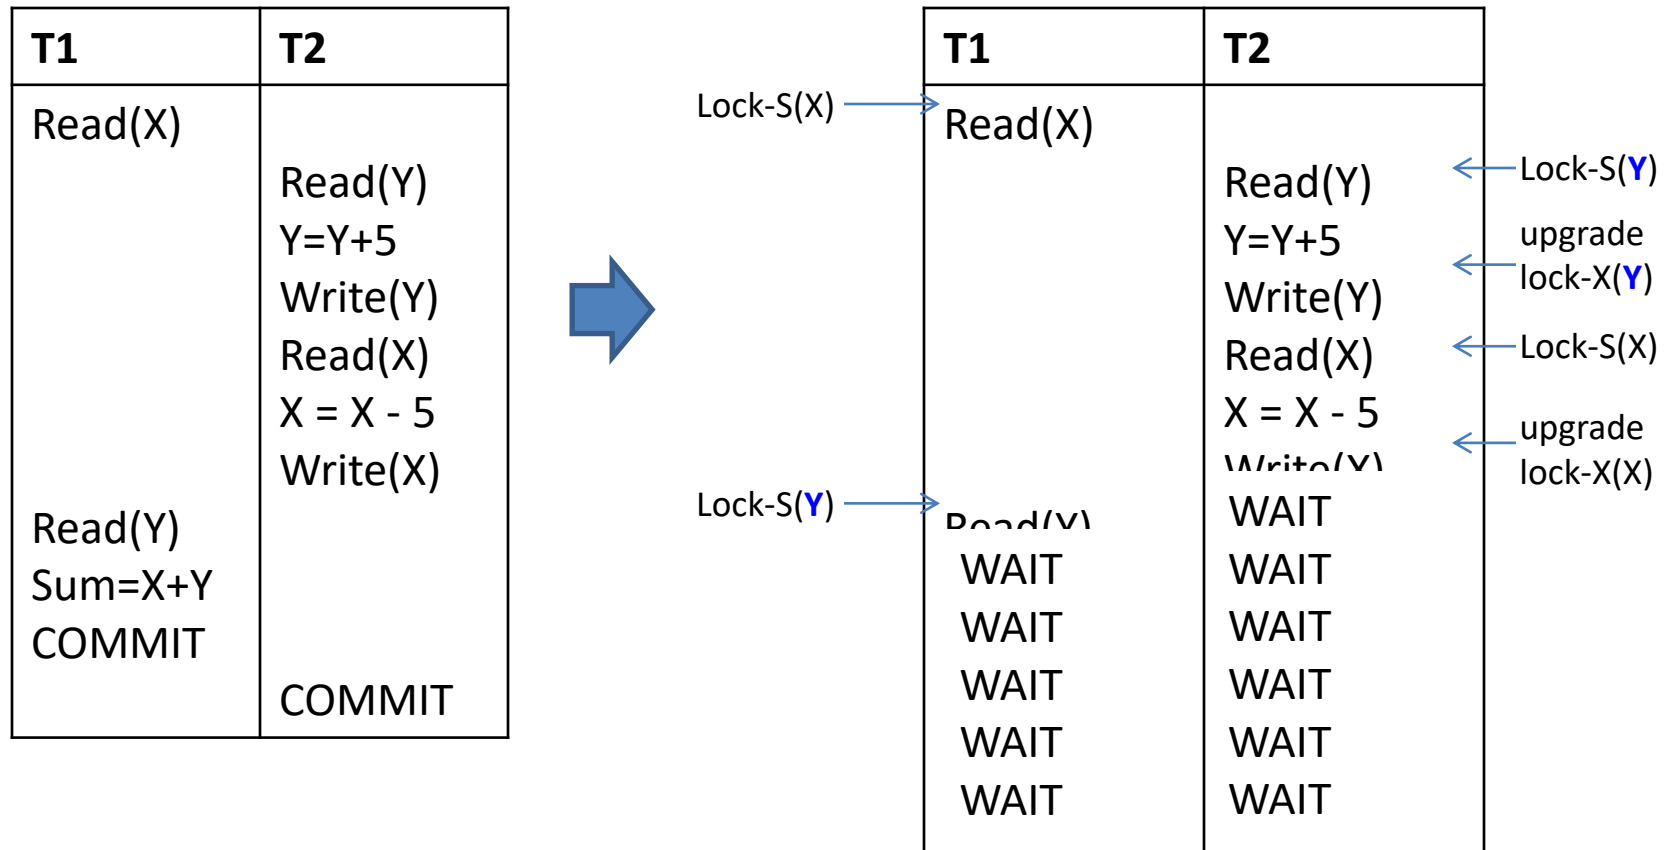


Uncommitted Update – problem solved.

Strict 2PL - Inconsistent Analysis

● Inconsistent Analysis

Lock	X	S
X	False	False
S	False	True



Strict 2PL prevent Inconsistent Analysis , and not allow to continue.
That's important: No inconsistency in database. Good!

This lecture


Overview

- Concurrency Protocols
- Involves two parts:
 - Locking Protocol
 - **Deadlock Detection and Handling**



Deadlock

- A deadlock is an impasse that may result when two or more transactions are waiting for locks to be released which are held by each other.



Andy and Becky both waits for each other to release their locks, and do nothing...

	T1	T2
Lock-S(X) →	Read(X)	
Lock-X(X) →	Write(X)	
		Read(Y) ← Lock-S(Y)
		Write(Y) ← Lock-X(Y)
Lock-X(Y) →	Read(Y)	
	WAIT	Read(X) ← Lock-X(X)
	WAIT	WAIT
	WAIT	WAIT
	WAIT	WAIT
	WAIT	WAIT
	WAIT	WAIT

Deadlock Detection

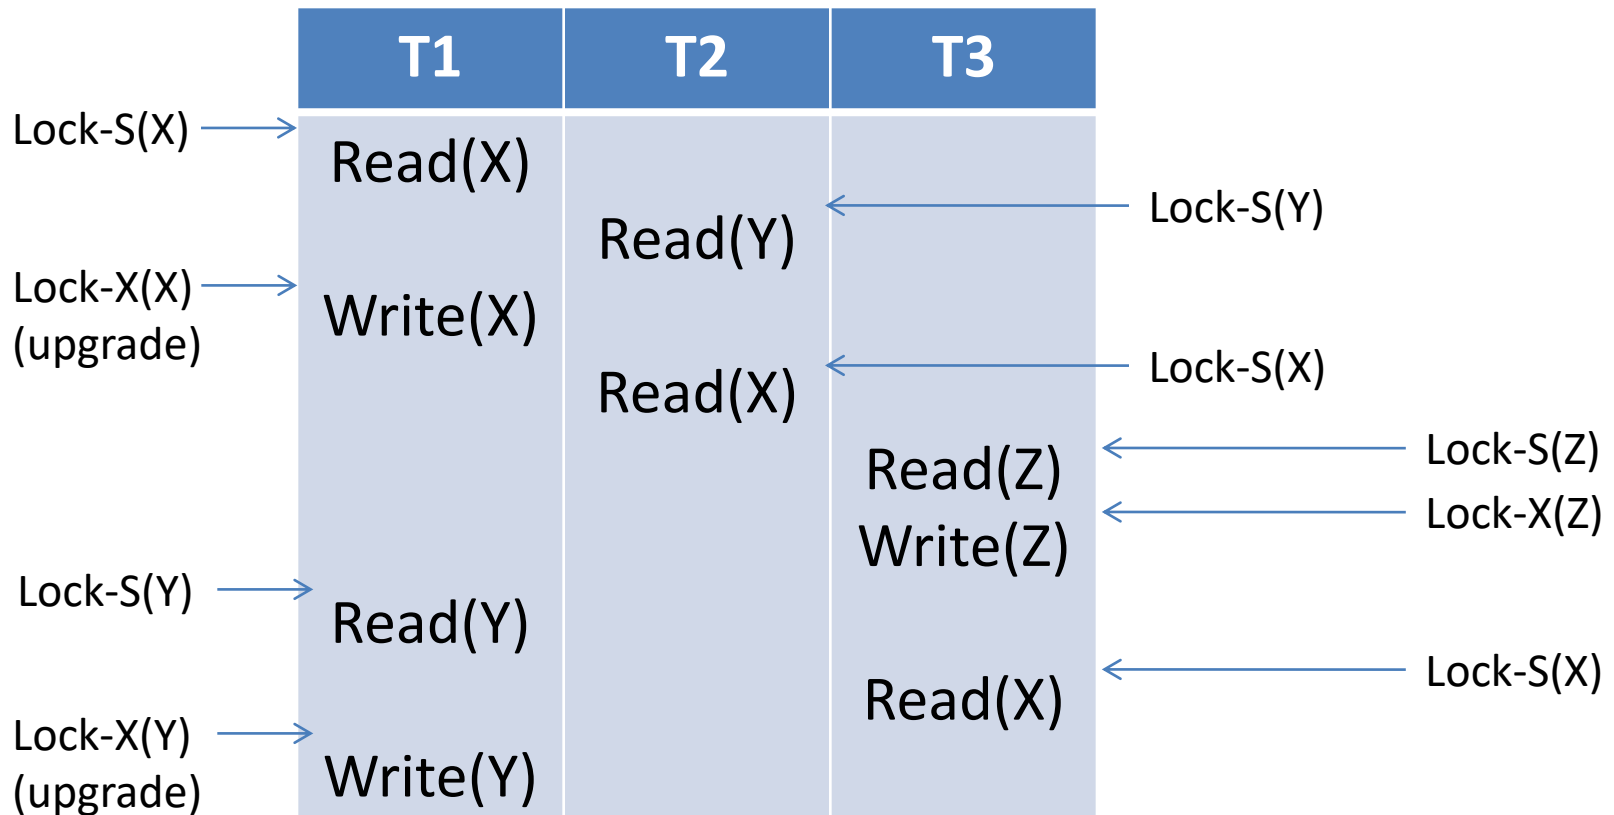
- Is there a deadlock?

T1	T2	T3
Read(X)	Read(Y) Read(X)	Read(Z) Write(Z) Read(X)
Write(X)		
Read(Y)		
Write(Y)		

	X	S
X	False	False
S	False	True

Deadlock Detection

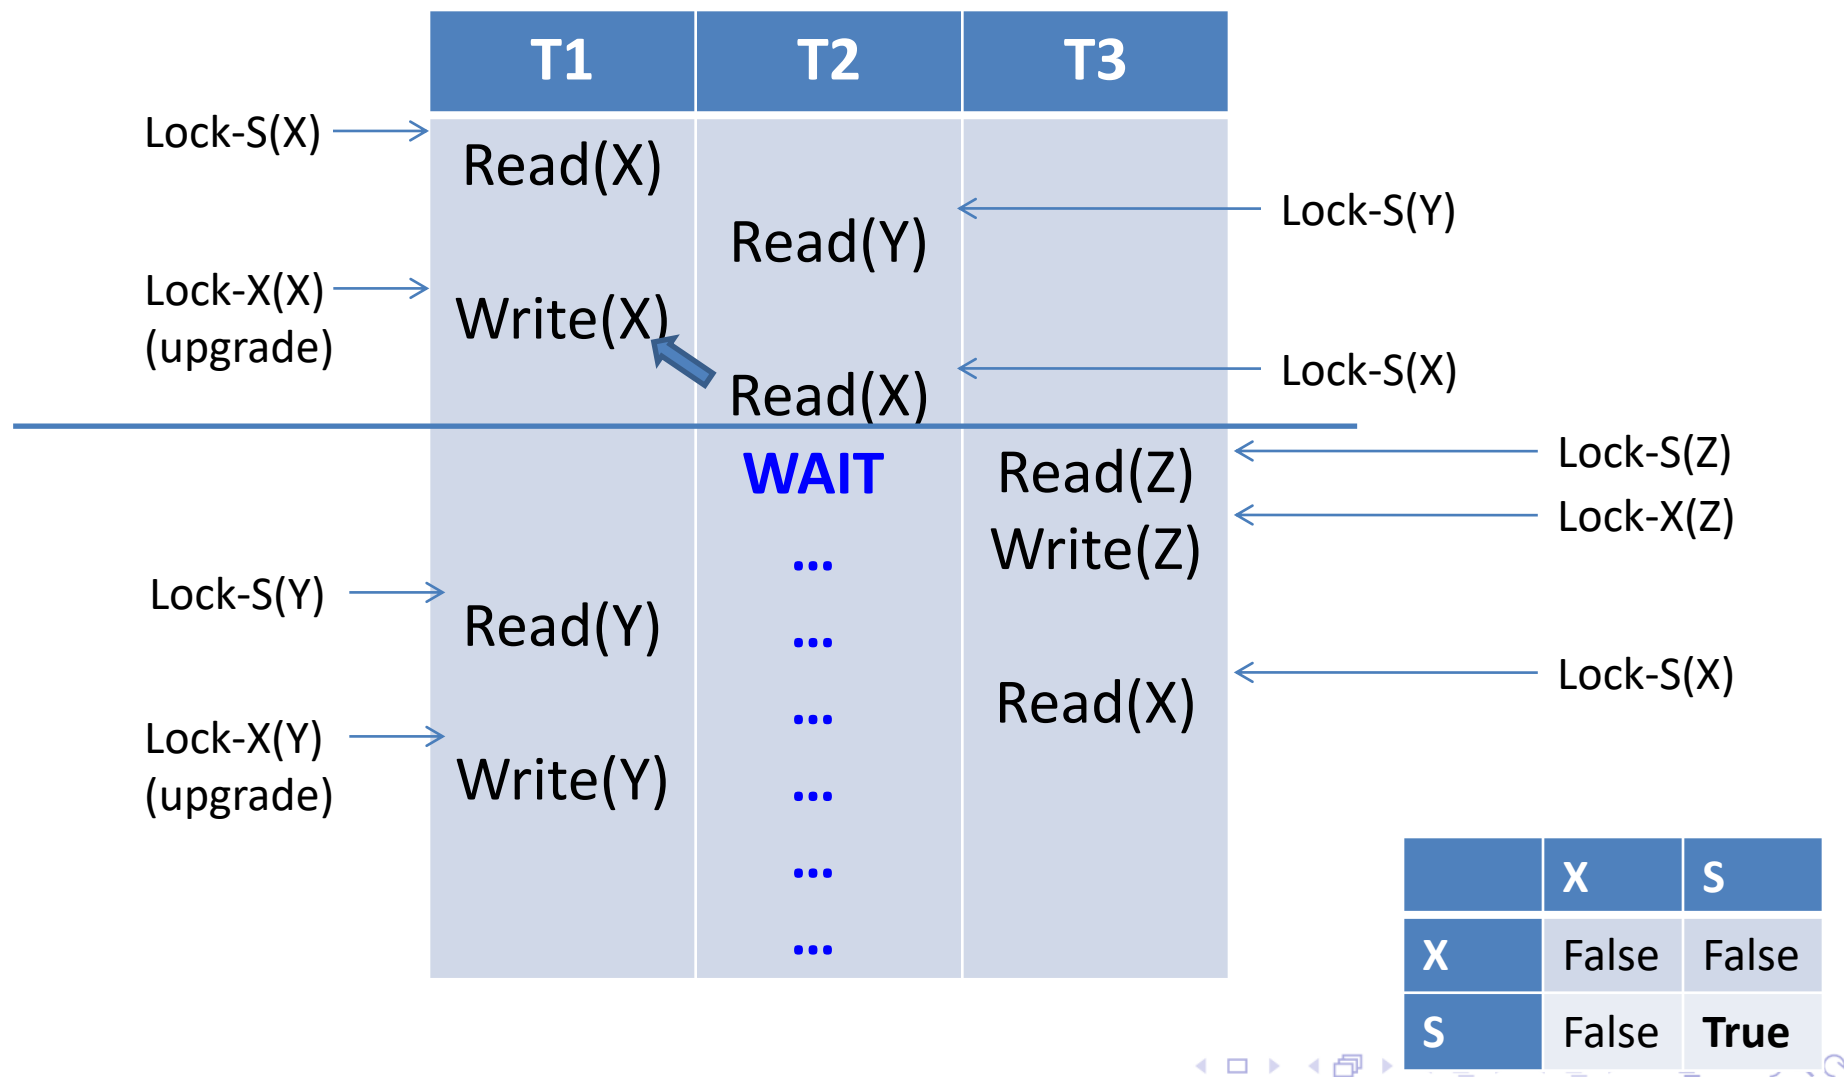
- Is there a deadlock? First write down the locks



	X	S
X	False	False
S	False	True

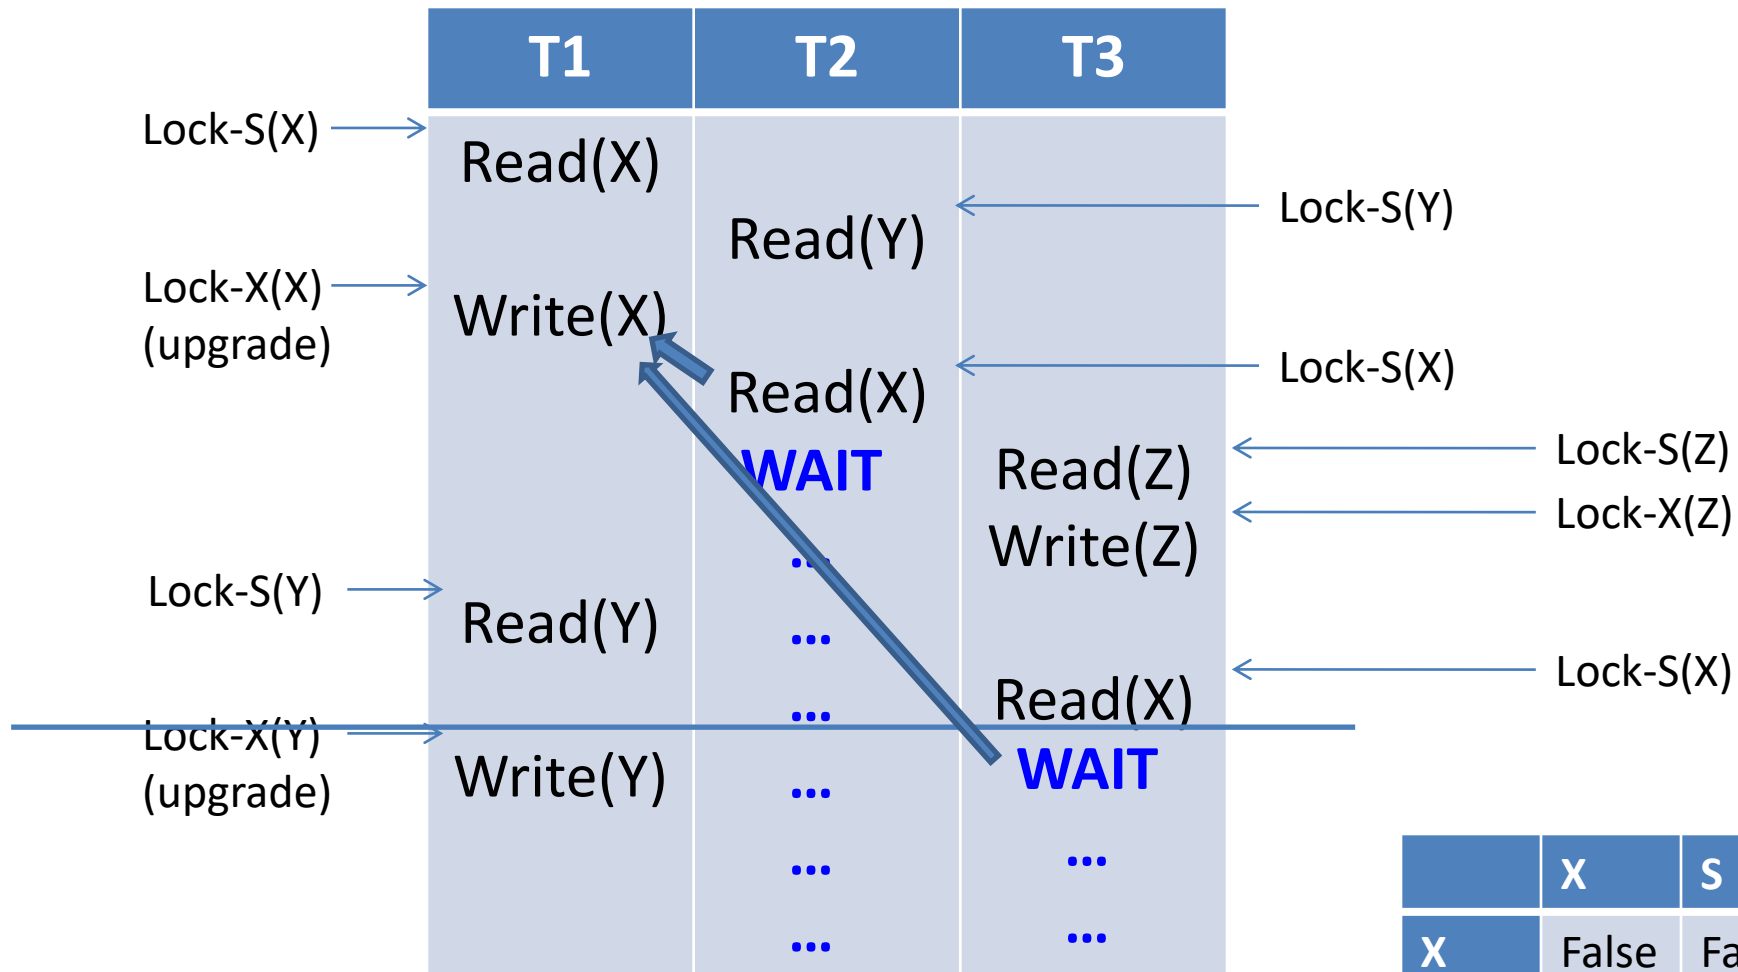
Deadlock Detection

- Is there a deadlock? Then, check locks above.



Deadlock Detection

- Is there a deadlock?



	X	S
X	False	False
S	False	True

- Is there a deadlock?

Wait-for graph

- Computer needs an algorithm
- Given a schedule, we can *detect* deadlocks which will happen in this schedule using a *wait-for graph* (WFG).
- Each transaction is a vertex
- Edge from *T2* to *T1* if
 - T1 Lock-S(X) then T2 tries to Lock-X(X) it
 - T1 Lock-X(X) then T2 tries to Lock-S(X) it
 - T1 Lock-X(X) then T2 tries to Lock-X(X) it

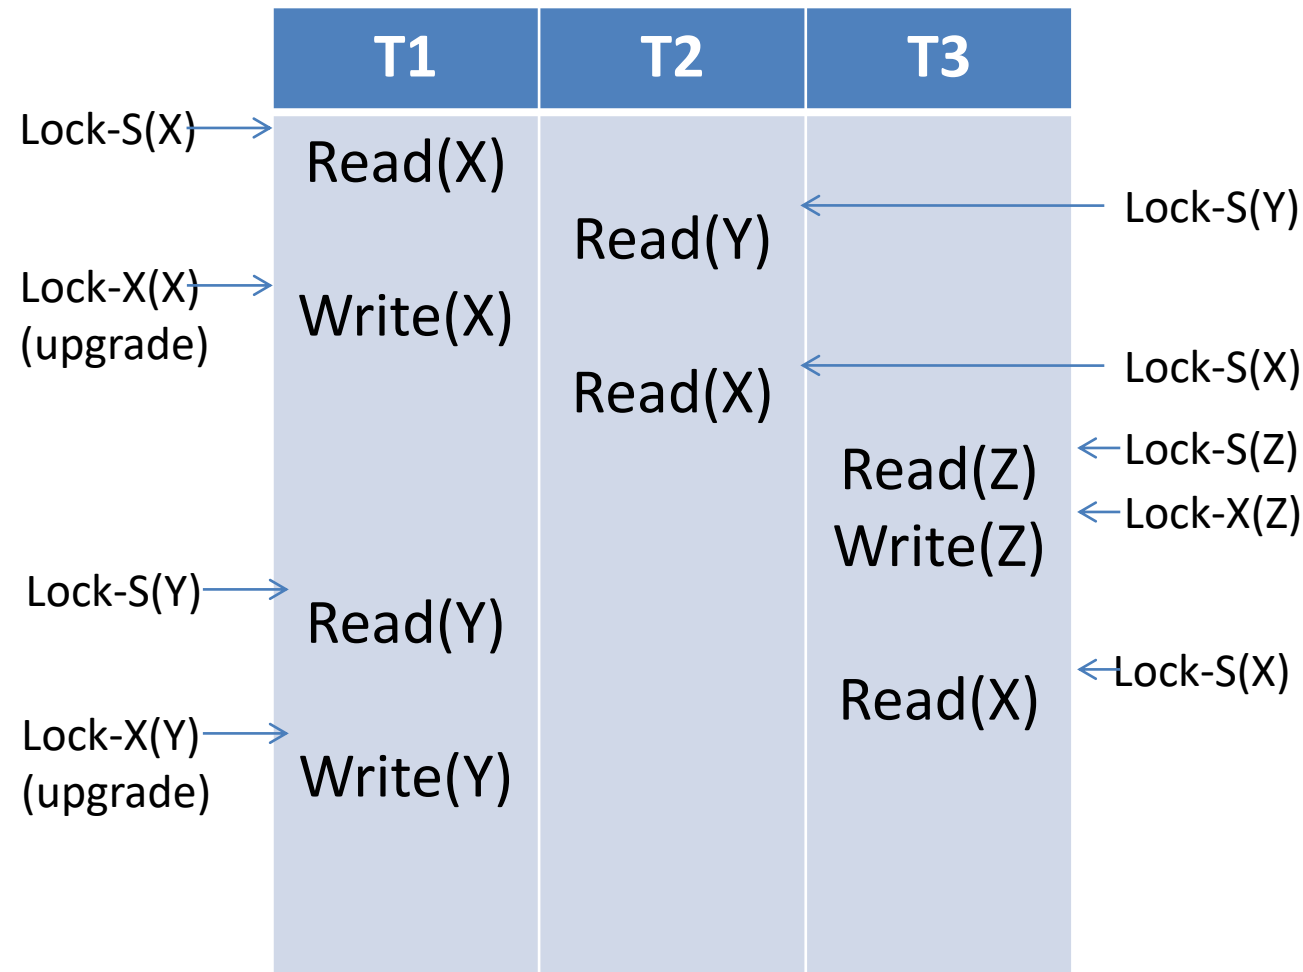
Locks operations:

- Look above
- Look at another transaction

Lock	X	S
X	False	False
S	False	True

Example – Wait For Graph

• Deadlock?




← Lock-S(Y)

← Lock-S(X)

← Lock-S(Z)

← Lock-X(Z)

← Lock-S(X)




Wait For Graph

Example – Wait For Graph

• Deadlock?


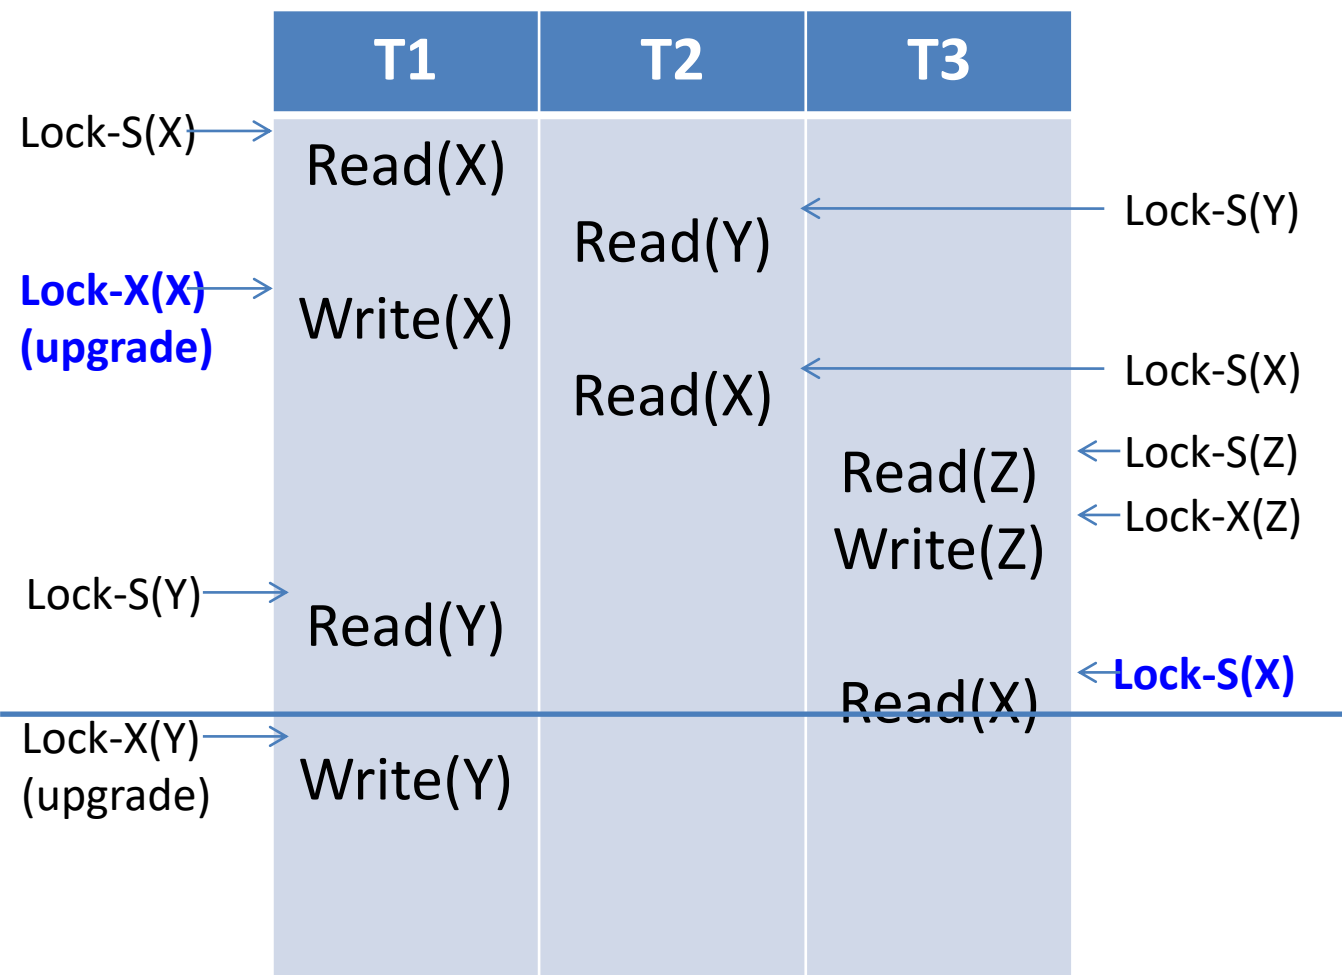
	T1	T2	T3
Lock-S(X) →	Read(X)		
		Read(Y)	
Lock-X(X) (upgrade) →	Write(X)		
		Read(X)	
<hr/>			
			Read(Z)
			Write(Z)
Lock-S(Y) →	Read(Y)		
			Read(X)
Lock-X(Y) (upgrade) →	Write(Y)		



Wait For Graph

Example – Wait For Graph


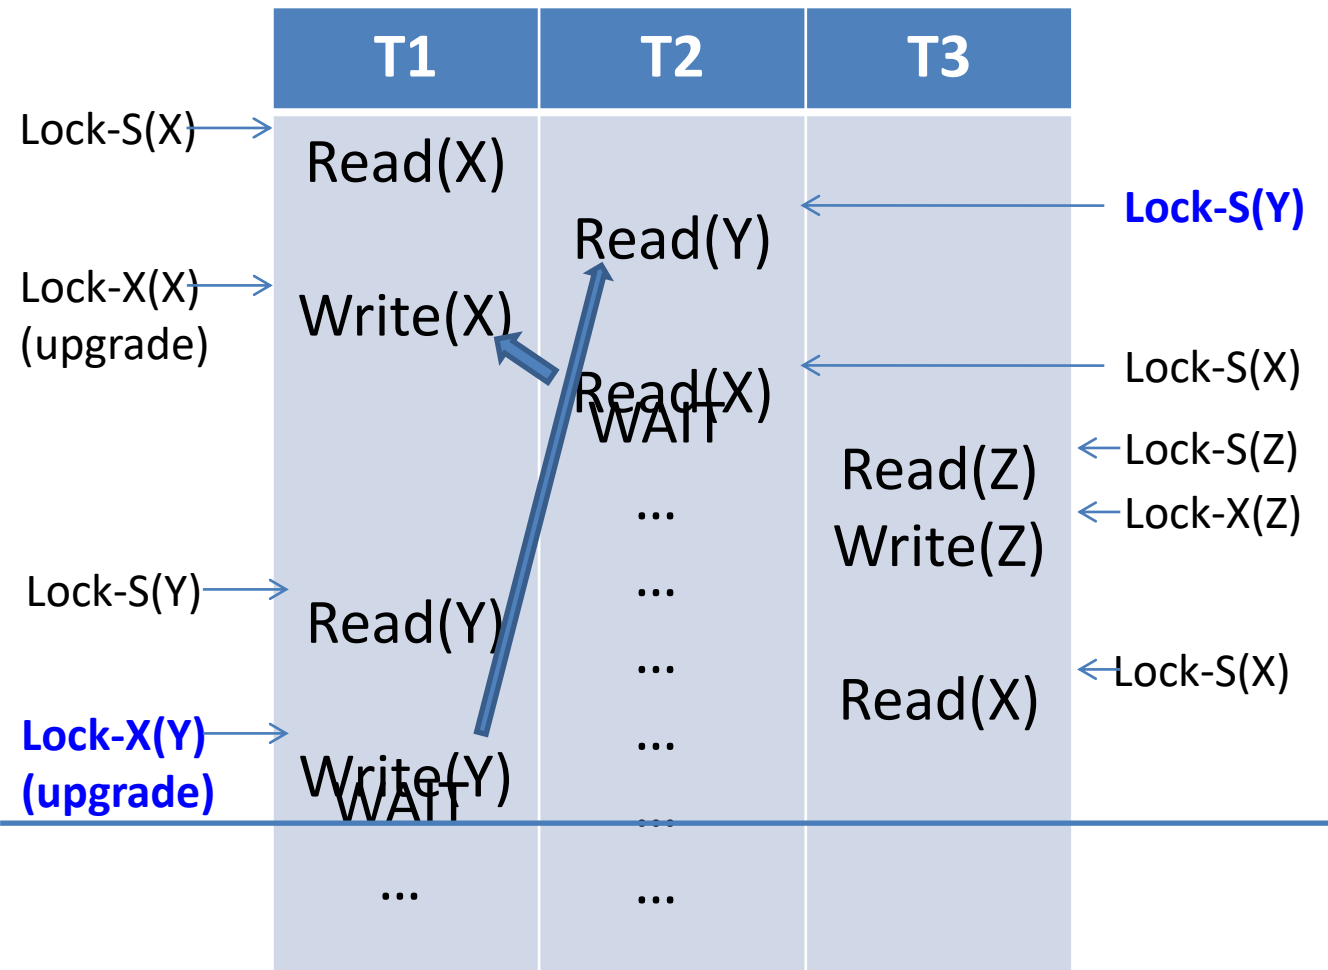
• Deadlock?



Wait For Graph

Example – Wait For Graph

• Deadlock?



Wait For Graph

Run DFS, Cycle detected : **Deadlock!**

Deadlock Recovery



- Deadlock is **less of a** problem than an **inconsistent** DB
- Most DBMSs (include Oracle, MySQL)
 - Uses **Strict 2PL** for locking
 - **Detect** deadlocks with a wait-for graph
 - **Recovery** : choose a single transaction as a 'victim' to rollback and restart
 - Which transaction?
 - has been running the longest/shortest?
 - have made the most/least updates?
 - have the most/least updates still to make?

